

Regression Models Course Project

August 20, 2020

You work for Motor Trend, a magazine about the automobile industry. Looking at a data set of a collection of cars, they are interested in exploring the relationship between a set of variables and miles per gallon (MPG) (outcome). They are particularly interested in the following two questions:

“Is an automatic or manual transmission better for MPG”

“Quantify the MPG difference between automatic and manual transmissions”

Take the mtcars data set and write up an analysis to answer their question using regression models and exploratory data analyses.

A data frame with 32 observations on 11 (numeric) variables.

mpg - Miles/(US) gallon

cyl - Number of cylinders

disp - Displacement (cu.in.)

hp - Gross horsepower

drat - Rear axle ratio

wt - Weight (1000 lbs)

qsec - 1/4 mile time

vs - Engine (0 = V-shaped, 1 = straight)

am - Transmission (0 = automatic, 1 = manual)

gear - Number of forward gears

carb - Number of carburetors

0.0.1 Import Libraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn

import xgboost as xgb
from xgboost import XGBClassifier, XGBRegressor
```

```

from xgboost import to_graphviz, plot_importance

#from sklearn.experimental import enable_hist_gradient_boosting
#from sklearn.ensemble import _hist_gradient_boosting
#from sklearn.ensemble import HistGradientBoostingRegressor,
↳HistGradientBoostingRegressor

%matplotlib inline
sns.set_style('dark')
sns.set(font_scale=1.2)

from sklearn.model_selection import cross_val_score, train_test_split,
↳GridSearchCV, RandomizedSearchCV
from sklearn.preprocessing import LabelEncoder, StandardScaler, MinMaxScaler,
↳OneHotEncoder
from sklearn.metrics import confusion_matrix, classification_report,
↳mean_absolute_error, mean_squared_error, r2_score
from sklearn.metrics import plot_confusion_matrix, plot_precision_recall_curve,
↳plot_roc_curve, accuracy_score
from sklearn.metrics import auc, f1_score, precision_score, recall_score,
↳roc_auc_score

import warnings
warnings.filterwarnings('ignore')

import pickle
from pickle import dump, load

np.random.seed(0)

#from pycaret.classification import *
#from pycaret.clustering import *
from pycaret.regression import *

pd.set_option('display.max_columns',100)
#pd.set_option('display.max_rows',100)
pd.set_option('display.width', 1000)

```

0.0.2 Data Exploration and Analysis

```
[2]: df = pd.read_csv("mtcars.csv")
```

```
[3]: df
```

```
[3]:
```

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear
carb											

0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4
4											
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4
4											
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4
1											
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3
1											
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3
2											
5	Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3
1											
6	Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3
4											
7	Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4
2											
8	Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4
2											
9	Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4
4											
10	Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4
4											
11	Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3
3											
12	Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3
3											
13	Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3
3											
14	Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3
4											
15	Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3
4											
16	Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3
4											
17	Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4
1											
18	Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4
2											
19	Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4
1											
20	Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3
1											
21	Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3
2											
22	AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3
2											
23	Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3

```

4
24      Pontiac Firebird  19.2    8  400.0  175  3.08  3.845  17.05  0  0    3
2
25          Fiat X1-9  27.3    4   79.0   66  4.08  1.935  18.90  1  1    4
1
26      Porsche 914-2  26.0    4  120.3   91  4.43  2.140  16.70  0  1    5
2
27      Lotus Europa  30.4    4   95.1  113  3.77  1.513  16.90  1  1    5
2
28      Ford Pantera L  15.8    8  351.0  264  4.22  3.170  14.50  0  1    5
4
29      Ferrari Dino  19.7    6  145.0  175  3.62  2.770  15.50  0  1    5
6
30      Maserati Bora  15.0    8  301.0  335  3.54  3.570  14.60  0  1    5
8
31      Volvo 142E  21.4    4  121.0  109  4.11  2.780  18.60  1  1    4
2

```

```
[4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32 entries, 0 to 31
Data columns (total 12 columns):
#   Column  Non-Null Count  Dtype
---  -
0   model   32 non-null      object
1   mpg     32 non-null      float64
2   cyl     32 non-null      int64
3   disp    32 non-null      float64
4   hp      32 non-null      int64
5   drat    32 non-null      float64
6   wt      32 non-null      float64
7   qsec    32 non-null      float64
8   vs      32 non-null      int64
9   am      32 non-null      int64
10  gear    32 non-null      int64
11  carb    32 non-null      int64
dtypes: float64(5), int64(6), object(1)
memory usage: 3.1+ KB

```

```
[5]: df.describe(include='all')
```

```

[5]:
      model      mpg      cyl      disp      hp      drat
wt      qsec      vs      am      gear      carb
count      32  32.000000  32.000000  32.000000  32.000000  32.000000
32.000000  32.000000  32.000000  32.000000  32.000000  32.0000
unique      32      NaN      NaN      NaN      NaN      NaN

```

NaN	NaN	NaN	NaN	NaN	NaN	
top	Fiat X1-9	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN	NaN	NaN	NaN	
freq	1	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN	NaN	NaN	NaN	
mean	NaN	20.090625	6.187500	230.721875	146.687500	3.596563
3.217250	17.848750	0.437500	0.406250	3.687500	2.8125	
std	NaN	6.026948	1.785922	123.938694	68.562868	0.534679
0.978457	1.786943	0.504016	0.498991	0.737804	1.6152	
min	NaN	10.400000	4.000000	71.100000	52.000000	2.760000
1.513000	14.500000	0.000000	0.000000	3.000000	1.0000	
25%	NaN	15.425000	4.000000	120.825000	96.500000	3.080000
2.581250	16.892500	0.000000	0.000000	3.000000	2.0000	
50%	NaN	19.200000	6.000000	196.300000	123.000000	3.695000
3.325000	17.710000	0.000000	0.000000	4.000000	2.0000	
75%	NaN	22.800000	8.000000	326.000000	180.000000	3.920000
3.610000	18.900000	1.000000	1.000000	4.000000	4.0000	
max	NaN	33.900000	8.000000	472.000000	335.000000	4.930000
5.424000	22.900000	1.000000	1.000000	5.000000	8.0000	

```
[6]: df.shape
```

```
[6]: (32, 12)
```


```
[7]: df.columns
```

```
[7]: Index(['model', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am',
        'gear', 'carb'], dtype='object')
```


0.0.3 Data Visualization

0.0.4 Univariate Data Exploration

```
[8]: df.hist(bins=50, figsize=(20,15))
plt.tight_layout()
plt.show()
```



```
[9]: df.boxplot(figsize=(20,15))
plt.tight_layout()
plt.show()
```



0.0.5 Bivariate Data Exploration

```
[10]: sns.jointplot(x='disp', y='mpg',data=df, kind='scatter')

sns.jointplot(x='hp', y='mpg',data=df, kind='scatter')

sns.jointplot(x='drat', y='mpg',data=df, kind='scatter')

sns.jointplot(x='wt', y='mpg',data=df, kind='scatter')

sns.jointplot(x='qsec', y='mpg',data=df, kind='scatter')


sns.jointplot(x='vs', y='mpg',data=df, kind='scatter')


sns.jointplot(x='am', y='mpg',data=df, kind='scatter')


# sns.jointplot(x='', y='',data=df, kind='reg')


# sns.jointplot(x='', y='',data=df, kind='reg')
```


```
plt.show()
```


















```
[11]: sns.barplot(x='am', y='mpg',data=df, ci=None);
```



Result: Manual gear gives higher MPG outcome by 10

```
[12]: df.groupby("am").mean()
```

```
[12]:
```

	mpg	cyl	disp	hp	drat	wt	qsec
vs	gear	carb					
am							
0	17.147368	6.947368	290.378947	160.263158	3.286316	3.768895	18.183158
	0.368421	3.210526	2.736842				
1	24.392308	5.076923	143.530769	126.846154	4.050000	2.411000	17.360000
	0.538462	4.384615	2.923077				

0.0.6 Correlation

```
[13]: df.corr()
```

```
[13]:
```

	mpg	cyl	disp	hp	drat	wt	qsec
vs	am	gear	carb				
mpg	1.000000	-0.852162	-0.847551	-0.776168	0.681172	-0.867659	0.418684
	0.664039	0.599832	0.480285	-0.550925			
cyl	-0.852162	1.000000	0.902033	0.832447	-0.699938	0.782496	-0.591242
	-0.810812	-0.522607	-0.492687	0.526988			
disp	-0.847551	0.902033	1.000000	0.790949	-0.710214	0.887980	-0.433698
	-0.710416	-0.591227	-0.555569	0.394977			

```


hp    -0.776168  0.832447  0.790949  1.000000 -0.448759  0.658748 -0.708223
-0.723097 -0.243204 -0.125704  0.749812
drat   0.681172 -0.699938 -0.710214 -0.448759  1.000000 -0.712441  0.091205
0.440278  0.712711  0.699610 -0.090790
wt    -0.867659  0.782496  0.887980  0.658748 -0.712441  1.000000 -0.174716
-0.554916 -0.692495 -0.583287  0.427606
qsec   0.418684 -0.591242 -0.433698 -0.708223  0.091205 -0.174716  1.000000
0.744535 -0.229861 -0.212682 -0.656249
vs     0.664039 -0.810812 -0.710416 -0.723097  0.440278 -0.554916  0.744535
1.000000  0.168345  0.206023 -0.569607
am     0.599832 -0.522607 -0.591227 -0.243204  0.712711 -0.692495 -0.229861
0.168345  1.000000  0.794059  0.057534
gear   0.480285 -0.492687 -0.555569 -0.125704  0.699610 -0.583287 -0.212682
0.206023  0.794059  1.000000  0.274073
carb  -0.550925  0.526988  0.394977  0.749812 -0.090790  0.427606 -0.656249
-0.569607  0.057534  0.274073  1.000000

```

```

[14]: plt.figure(figsize=(16,9))
sns.heatmap(df.corr(),cmap="coolwarm",annot=True,fmt='.2f',linewidths=2)
plt.show()


```



```

[15]: sns.pairplot(df)
plt.show()

```

0.0.7 Data Preprocessing

0.0.8 Treat Missing Values

```
[16]: df.isnull().sum()
```

```
[16]: model    0
      mpg      0
      cyl      0
      disp     0
      hp       0
      drat     0
      wt       0
```

```

qsec      0
vs         0
am         0
gear       0
carb       0
dtype: int64

```

0.0.9 Treat Duplicate Values

```
[17]: df.duplicated(keep='first').sum()
```

```
[17]: 0
```

0.0.10 Drop unwanted features

```
[18]: df.columns
```

```
[18]: Index(['model', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am',
'gear', 'carb'], dtype='object')
```

```
[19]: df.drop(['model'],axis=1,inplace=True)
```

```
[20]: df
```

```
[20]:
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
5	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
6	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
7	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
8	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
9	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
10	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
11	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
12	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
13	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
14	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
15	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
16	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
17	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
18	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
19	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
20	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1

21	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
22	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
23	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
24	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
25	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
26	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
27	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
28	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
29	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
30	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
31	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

0.0.11 Create and save processed dataset

```
[21]: #df.to_csv("carstrain.csv",index=False)
```

0.0.12 Model Training

```
[22]: df.columns
```

```
[22]: Index(['mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am', 'gear',  
          'carb'], dtype='object')
```

0.0.13 Using PyCaret

```
[23]: exp_reg = setup(data = df, target = 'mpg', session_id=0, normalize=True,  
                    ↪categorical_features=['vs', 'am', 'gear', 'carb'])
```

Setup Succesfully Completed.

<pandas.io.formats.style.Styler at 0x1c0ef41b400>

```
[24]: compare_models()
```

<pandas.io.formats.style.Styler at 0x1c0ec264640>

```
[24]: BayesianRidge(alpha_1=1e-06, alpha_2=1e-06, alpha_init=None,  
                  compute_score=False, copy_X=True, fit_intercept=True,  
                  lambda_1=1e-06, lambda_2=1e-06, lambda_init=None, n_iter=300,  
                  normalize=False, tol=0.001, verbose=False)
```